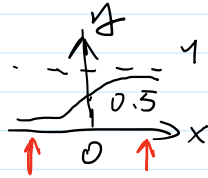


Week2 - Logistic Regression with Gradient Descent

Donnerstag, 16. Juli 2020 19:32

Logistic Regression:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



$$\left. \begin{array}{l} x \rightarrow -\infty \quad y=0 \\ x \rightarrow +\infty \quad y=1 \end{array} \right\}$$

Binary Classification problem

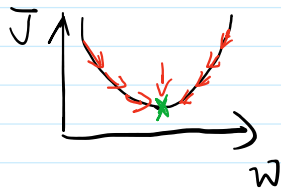
$$\hat{y} = \sigma(w^T x + b) \quad \text{with } w^T \in \mathbb{R}^n \quad \hat{y} \text{ predicted value}$$

$$b \in \mathbb{R} \quad y \text{ ground true}$$

$$\text{loss / error} \rightarrow \mathcal{L}(\hat{y}, y) = - (y \log(\hat{y}) + (1-y) \log(1-\hat{y}))$$

$$\text{cost function} \rightarrow J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}_i, y_i)$$

Gradient descent:



Repeat {

$$w := w - \alpha \underbrace{\frac{dJ(w)}{dw}}_{\text{gradient}}$$

}

$\alpha \rightarrow$ learning rate

when learning rate too large, fast but may never converge;
when learning rate too small, will converge however slow.

Vectorization:

Vectorization can greatly increase the process and simplify the algorithms.

$$X = \begin{bmatrix} \underbrace{x^{(1)}}_n & \underbrace{x^{(2)}}_n & \underbrace{x^{(3)}}_n & \dots & \underbrace{x^{(m)}}_n \end{bmatrix} \quad X \in \mathbb{R}^{n \times m}$$

$$z = [z^{(1)} \quad z^{(2)} \quad \dots \quad z^{(m)}] = \underbrace{w^T X}_{1 \times m} + \underbrace{[b \quad b \quad b \dots b]}_{1 \times m}$$

$$z = n \cdot \text{dot}(w^T, X) + b \rightarrow \text{broadcasting to } 1 \times m$$

$$A = [a^{(1)} \quad a^{(2)} \quad a^{(3)} \quad \dots \quad a^{(m)}] = \sigma(z) \quad z \Rightarrow (1 \times m)$$

$$Y = [y^{(1)} \quad y^{(2)} \quad \dots \quad y^{(m)}]$$

$$dz = A - Y = [a^{(1)} - y^{(1)} \quad a^{(2)} - y^{(2)} \quad \dots \quad a^{(m)} - y^{(m)}] \rightarrow 1 \times m$$

$$\begin{cases} dw = \frac{1}{m} \times dz^T & x \rightarrow (n \times m) \quad dz^T \rightarrow (m \times 1) \Rightarrow (n \times 1) \\ db = \frac{1}{m} \cdot \text{np.sum}(dz) & (1 \times 1) \end{cases}$$

Take out from programming assignment:

----- Get familiar with Numpy -----

Sigmoid function : $x \rightarrow \text{numpy array}$

$$s = \text{sigmoid}(x)$$

$$= 1 / (1 + \text{np.exp}(-x))$$

Sigmoid gradient : $s'(x) = s(x)(1 - s(x))$

An image is represented by a 3D array of shape (length, height, depth = 3); Yet, when you read an image as the input of an algorithm you convert it to a vector of shape (length*height*3, 1), so to speak a column vector, 1D vector.



```
v = image.reshape((image.shape[2]*image.shape[0]*image.shape[1], 1))
```

with image -- a numpy array of shape (length, height, depth)

Normalizing input data helps algorithms to converge faster: dividing each row vector of x by its norm.

```
x_norm = np.linalg.norm(x, ord = 2, axis = 1, keepdims = True)
x = x / x_norm
```

Matrix Operation:

dot product: `np.dot(x1,x2)` result is scalar

outer product: `np.outer(x1,x2)`

Elementwise Multiplication: `np.multiply(x1,x2)` or `x1 * x2`

Matrix Multiplication: `x1 @ x2`

----- Logistic Regression with a Neural Network mindset -----

The first step to a classifier is commonly data preprocessing.

1. Check the dimensions and shapes of the datasets, including training dataset, test dataset, etc. #
2. The dataset should normally has a shape of ((num_of_pixel * num_of_pixel * 3, 1))
3. In case of images, reshape the datasets into vector of size (height * width * 3, 1) may be helpful
 -> `X_flatten = X.reshape(X.shape[0], -1).T`
 Or `X_flatten = X.reshape(X.shape[0] * X.shape[1] * 3, 1)`
4. To represent color images, the red, green and blue channels (RGB) must be specified for each pixel, and so the pixel value is actually a vector of three numbers ranging from 0 to 255. One common preprocessing step is to center and standardize the dataset, meaning:
 - a. Subtracting the mean of the whole numpy array from each example, and then divide each example by the

- standard deviation of the whole numpy array.
- b. For image input, subtracting each pixel value by 255.

After preprocessing of the datasets, the logistic regression itself is implemented in:

$$z^{(i)} = w^T x^{(i)} + b$$

$$\hat{y}^{(i)} = a^{(i)} = \text{sigmoid}(z^{(i)})$$

$$\mathcal{L}(a^{(i)}, y^{(i)}) = -y^{(i)} \log(a^{(i)}) - (1 - y^{(i)}) \log(1 - a^{(i)})$$

elementwise

$$J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(a^{(i)}, y^{(i)})$$

General structure of a NN algorithms:

- Define the model structure (such as number of input features, hidden layers, etc.)
- Initialize the model's parameters
- Loop(vectorized):
 - Calculate current loss (forward propagation)
 - Calculate current gradient (backward propagation)
 - Update parameters (gradient descent)

The previous steps can be implemented in different functions and then called in one model()

Steps:

1. initialize parameters (w, b)
 - > `w = np.zeros((dim,1)); b = 0`
2. Forward propagation:
 - a. sigmoid() or other transfer function
 - b. activation `A = sigmoid(w*X + b)`
 - c. cost function `L`
3. Back propagation:
 - a. `dw = dL/dw`
 - b. `db = dL/db`
4. Optimizing (w, b):
 - a. for `i` in `range(number of iterations)`
 - i. `w = w - learning_rate * dw`
 - ii. `b = b - learning_rate * db`

```
def optimize(w, b, X, Y, num_iterations, learning_rate):
```

Gradient descent

Returns:

params -- dictionary containing the weights w and bias b
 grads -- dictionary containing the gradients of the weights and bias with respect to the cost function
 costs -- list of all the costs computed during the optimization, this will be used to plot the learning curve.

...

Cost = []

for i in range(num_iterations):

Cost and gradient calculation
 grads, cost = propagate(w, b, X, Y)

derivatives from grads
 dw = grads["dw"]
 db = grads["db"]

update
 w = w - learning_rate * dw
 b = b - learning_rate * db

Record the costs every 100 iterations and print on screen
 if i % 100 == 0:
 costs.append(cost)
 print ("Cost after iteration %i: %f" % (i, cost))

params = {"w": w, "b": b}
 grads = {"dw": dw, "db": db}
 return params, grads, costs

5. Predict using optimized (w, b)
 - a. if `A = sigmoid(w*X + b) > 0.5`:
 - i. `Y_predict = 1`
 - b. else:
 - i. `y_predic = 0`
6. integrate all previous steps into a learning model

```
def model(X_train, Y_train, X_test, Y_test, num_iterations = 2000, learning_rate = 0.5):
```

```
def propagation(w, b, X, Y):
    ...
    Implement cost function and gradients

    Arguments:
    w -- weights, a numpy array of size (num_px * num_px * 3, 1)
    b -- bias, a scalar
    X -- data of size (num_px * num_px * 3, number of examples)
    Y -- true "label" vector (containing 0 if non-cat, 1 if cat) of size (1, number of examples)
    ...

    m = X.shape[1]; # number of examples

    # Forward Prog
    A = sigmoid(np.dot(w.T, X) + b) # (1*m)
    cost = -1/m * np.sum(Y * np.log(A) + (1-Y) * np.log(1-A))

    # Back Prog
    dw = 1/m * np.dot(X, (A-Y).T) # (n*1)
    db = 1/m * np.sum(A-Y) # (1*1)
    grads = {"dw": dw, "db": db}

    Return grads, cost
```

$$dw = \frac{1}{m} X dz^T$$

```
def predict(w, b, X):
    ...
    Predict using learned (w,b)
    ...

    m = X.shape[1]
    Y_prediction = np.zeros((1,m))

    m = X.shape[1]
    Y_prediction = np.zeros((1,m))
    w = w.reshape(X.shape[0], 1)

    for i in range(A.shape[1]):

        # Convert probabilities A[0,i] to actual predictions
        p[0,i]
        if A[:, i] > 0.5:
            Y_prediction[:, i] = 1
        else:
            Y_prediction[:, i] = 0

    Return grads, cost
```

```

# initialize parameters with zeros
w, b = initialize_with_zeros(X_train.shape[0])

# Gradient descent
parameters, grads, costs = optimize(w, b, X_train, Y_train, num_iterations,
learning_rate, print_cost)

# Retrieve parameters w and b from dictionary "parameters"
w = parameters["w"]
b = parameters["b"]

# Predict
Y_prediction_test = predict(w, b, X_test)
Y_prediction_train = predict(w, b, X_train)

# Print train/test Errors
print("train accuracy: {} %".format(100 - np.mean(np.abs(Y_prediction_train -
Y_train)) * 100))
print("test accuracy: {} %".format(100 - np.mean(np.abs(Y_prediction_test -
Y_test)) * 100))

d = {"costs": costs,
      "Y_prediction_test": Y_prediction_test,
      "Y_prediction_train" : Y_prediction_train,
      "w" : w,
      "b" : b,
      "learning_rate" : learning_rate,
      "num_iterations": num_iterations}

return d

```

7. By plot d["cost"] against iterations we get the learning curve
 - a. Different learning rates give different costs and thus different predictions results.
 - b. If the learning rate is too large (0.01), the cost may oscillate up and down. It may even diverge.
 - c. In deep learning recommended:
 - Choose the learning rate that better minimizes the cost function.
 - If your model overfits, use other techniques to reduce overfitting.

Note: This is the implementation of a logistic regression problem with gradient decent approach. The Forward / Back propagation involves Neural network thinking and will be looked into detail in the further assignments.